

✓ Import Libraries and Load data

```
import pandas as pd
import numpy as np
```

```
treatment = pd.read_json("/content/treatment_group_data_raw.json")
control = pd.read_json("/content/control_group_data_raw.json")
```

✓ Preprocessing and Feature Engineering

```
treatment['release_date'] = pd.to_datetime(treatment['release_date'])
treatment['release_year'] = treatment['release_date'].dt.year
treatment['release_quarter'] = treatment['release_date'].dt.quarter
```

```
control['release_date'] = pd.to_datetime(control['release_date'])
control['release_year'] = control['release_date'].dt.year
control['release_quarter'] = control['release_date'].dt.quarter
```

```
mcu_movie_ids = [429617, 299534, 299537, 363088, 299536, 284054, 284053, 315635,
                 283995, 284052, 271110, 102899, 99861, 118340, 100402, 76338,
                 68721, 24428, 1771, 10195, 10138, 1724, 1726]
```

```
treatment['MCU'] = treatment.apply(lambda x: 1 if x['movie_id'] in mcu_movie_ids else 0)
```

```
np.int64(76)
```

```
treatment.drop(['movie_id', 'imdb_id'], axis=1, inplace=True)
control.drop(['movie_id', 'imdb_id'], axis=1, inplace=True)
```

✓ Constructing Treatment and Control Groups

```
treatment['Treatment'] = 1
control['Treatment'] = 0
```


```
treatment_control = pd.concat([treatment,control],axis=0)
```

```
from tqdm import tqdm
tqdm.pandas()
```

```
treatment_control['Lead Role'] = treatment_control.progress_apply(lambda x: 1 if
treatment_control['Supporting Role'] = treatment_control.progress_apply(lambda x:
```

```
↔ 100%|██████████| 670/670 [00:00<00:00, 66293.55it/s]
100%|██████████| 670/670 [00:00<00:00, 91426.74it/s]
```

treatment_control



	title	budget	revenue	runtime	release_date	cast_order	actor_name
272	My Life Without Me	2500000	12300000	106	2003-03-07	3	Mark Ruffalo
274	View from the Top	30000000	19526014	87	2003-03-21	2	Mark Ruffalo
1	The Shape of Things	0	735992	96	2003-07-24	0	Paul Rudd
335	S.W.A.T.	80000000	207700000	117	2003-08-08	5	Jeremy Renner
169	Lost in Translation	4000000	119723856	102	2003-09-18	1	Scarlett Johansson
...
137	Blink Twice	20000000	46393906	102	2024-08-21	1	Channing Tatum
308	The Killer	30000000	318618	126	2024-08-22	2	Sam Worthington
270	The Order	20000000	1970445	116	2024-12-05	2	Tye Sheridan
44	A Complete Unknown	65000000	138003641	140	2024-12-18	1	Edward Norton
12	Sonic the Hedgehog 3	122000000	486018457	110	2024-12-19	0	Jim Carrey

670 rows × 33 columns

✦ Identifying MCU Entry and Actor Treatment Timing

```
def keep_first_one_only(df, group_col='actor_name', flag_col='MCU'):  
    # Create a column tracking cumulative sum of 1s per group  
    df['_cumsum'] = df.groupby(group_col)[flag_col].cumsum()  
  
    # Set flag to 0 if it's a 1 and it's not the first one  
    df["MCU Entry"] = df.apply(lambda row: 1 if row[flag_col] == 1 and row['_cumsum'] == 1 else 0, axis=1)  
  
    # Drop helper column  
    df.drop(columns=['_cumsum'], inplace=True)  
    return df  
  
treatment_control = keep_first_one_only(treatment_control)  
  
treatment_control['MCU Entry Year'] = treatment_control.apply(lambda x: 0 if not x['MCU Entry'] else x['Year'], axis=1)  
  
treatment_control['actor_name'].unique()  
  
⇒ array(['Mark Ruffalo', 'Paul Rudd', 'Jeremy Renner', 'Scarlett Johansson',  
        'Robert Downey Jr.', 'Chris Evans', 'Chris Pratt',  
        'Benedict Cumberbatch', 'Chadwick Boseman', 'Chris Hemsworth',  
        'Tom Hiddleston', 'Tom Holland', 'Matthew Goode', 'Ben Foster',  
        'Jim Carrey', 'Ryan Reynolds', 'Steve Carell', 'Charlize Theron',  
        'Edward Norton', 'Sam Worthington', 'Channing Tatum',  
        'Tye Sheridan', 'Taron Egerton', 'John David Washington'],  
        dtype=object)
```

```
treatment_control[treatment_control['MCU Entry']==1]
```



	title	budget	revenue	runtime	release_date	cast_order	actor_name
215	Iron Man	140000000	585174222	126	2008-04-30	0	Rob Downey
172	Iron Man 2	200000000	623933331	124	2010-04-28	3	Scar Johanss
311	Thor	150000000	449326618	115	2011-04-21	2	Th Hiddlest
284	Thor	150000000	449326618	115	2011-04-21	0	Ch Hemswo
330	Thor	150000000	449326618	115	2011-04-21	50	Jerem Renr
91	Captain America: The First Avenger	140000000	370569774	124	2011-07-22	0	Chris Eva
279	The Avengers	220000000	1518815515	143	2012-04-25	2	Mark Ruffa
68	Guardians of the Galaxy	170000000	772776600	121	2014-07-30	0	Chris Pr
30	Ant-Man	130000000	519311965	117	2015-07-14	0	Paul Ru
75	Captain America: Civil War	250000000	1155046416	147	2016-04-27	12	Tom Holla
128	Captain America: Civil War	250000000	1155046416	147	2016-04-27	7	Chadw Bosem
166	Doctor Strange	180000000	676343174	115	2016-10-25	0	Benec Cumberbat

12 rows x 34 columns

✓ Visualizing Actor Treatment Timeline

```
heatmap_df = treatment_control.groupby('actor_name').agg({'MCU Entry Year':"sum",
heatmap_df.head()
```



	MCU Entry Year	Treatment
actor_name		
Ben Foster	0	0
Benedict Cumberbatch	2016	1
Chadwick Boseman	2016	1
Channing Tatum	0	0
Charlize Theron	0	0

✓ Staggered DiD Diagram

```
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.patches as mpatches

# --- Step 1: Setup base matrix ---
years = list(range(2003, 2025))
actors = heatmap_df.index.tolist()

# Create full grid: rows = actors, columns = years
status_matrix = pd.DataFrame(index=actors, columns=years)

# --- Step 2: Fill in treatment status ---
for actor in actors:
    entry_year = heatmap_df.loc[actor, 'MCU Entry Year']
    treated = heatmap_df.loc[actor, 'Treatment']

    for year in years:
        if treated == 0:
            status_matrix.loc[actor, year] = 0 # control
        elif year < entry_year:
```

```

        status_matrix.loc[actor, year] = 1 # treated (pre)
    else:
        status_matrix.loc[actor, year] = 2 # treated (post)

status_matrix = status_matrix.astype(float) # convert to numeric for heatmap

# --- Step 3: Sort actors by entry year ---
entry_order = heatmap_df[heatmap_df['Treatment'] == 1].sort_values('MCU Entry Year')
control_order = heatmap_df[heatmap_df['Treatment'] == 0].index.tolist()
status_matrix = status_matrix.loc[entry_order + control_order]

# --- Step 4: Plot heatmap ---
plt.figure(figsize=(14, 10))
cmap = sns.color_palette(["#cbd5e8", "#8da0cb", "#4c578a"]) # light -> dark


sns.heatmap(
    status_matrix,
    cmap=cmap,
    linewidths=0.2,
    linecolor='gray',
    cbar=False
)


plt.xlabel("Year")
plt.ylabel("Actor")

# Legend
legend_patches = [
    mpatches.Patch(color="#cbd5e8", label='Control (No MCU)'),
    mpatches.Patch(color="#8da0cb", label='Treated (Pre-MCU)'),
    mpatches.Patch(color="#4c578a", label='Treated (Post-MCU)')
]
plt.legend(handles=legend_patches, bbox_to_anchor=(1.02, 1), loc='upper left')

plt.tight_layout()
plt.show()

```






```
treatment_control.columns
```

```
⇒ Index(['title', 'budget', 'revenue', 'runtime', 'release_date', 'cast_order',
        'actor_name', 'imdb_votes', 'Internet Movie Database',
        'Rotten Tomatoes', 'Metacritic', 'opening_weekend', 'Action',
        'Adventure', 'Comedy', 'Drama', 'Family', 'Fantasy', 'Horror',
        'Music',
        'Mystery', 'Romance', 'Science Fiction', 'Thriller', 'War', 'Western',
        'MCU', 'MCU Entry', 'release_year', 'release_quarter', 'Treatment',
        'Lead Role', 'Supporting Role', 'MCU Entry Year'],
        dtype='object')
```

✓ Preliminary Data Analysis

```
import plotly.express as px
```

```
# ✓ Step 1: Choose the metric
```

```
df = treatment_control.copy()
```

```
#df = df[df['Treatment']==1]
```

```
df['profit'] = df['revenue'] - df['budget']
```

```
# ✓ Step 2: Filter data for years and valid entries
```

```
df = df[df['release_year'].between(2003, 2024)]
```

```
df = df.dropna(subset=['Treatment', 'release_year', 'budget', 'revenue'])
```

```
# Optional: limit to lead roles only
```

```
# df = df[df['Lead Role'] == 1]
```

```
# ✓ Step 3: Aggregate metric per actor per year
```

```
agg = df.groupby(['Treatment', 'release_year']).agg({
```

```
    'budget': 'mean',
```

```
    'revenue': 'mean',
```

```
    'opening_weekend': 'mean'
```

```
}).reset_index()
```

```
# ✓ Step 4: Plot with Plotly
```

```
fig = px.line(
```

```
    agg,
```

```
    x='release_year',
```

```
    y='opening_weekend', # Change to 'budget' or 'revenue' if needed
```

```
    color='Treatment',
```

```
    markers=True,
```

```

title='Profit per Treatment Group over Time (2003–2024)',
labels={'release_year': 'Year', 'opening_weekend': 'Opening Weekend Performance'},
hover_name='Treatment'
)


fig.update_layout(
    xaxis=dict(tickmode='linear', dtick=1),
    hovermode='x unified',
    legend_title='Treatment'
)

fig.show()

```



Profit per Treatment Group over Time (2003–2024)



```

treatment_control['opening_weekend_log'] = np.log(treatment_control["opening_weekend_log"])

```

```
import statsmodels.formula.api as smf
```

```
# 1) TWFE regression
twfe = smf.ols('opening_weekend_log ~ Treatment + C(actor_name) + C(release_year)
print("TWFE treat coef:", twfe.params['Treatment'], "p=", twfe.pvalues['Treatment
```

```
↔ TWFE treat coef: 1.9147393638815493 p= 0.0035044674802845077
```

```
treatment_control.T
```

↔		272	274	1	335	169	250
	title	My Life Without Me	View from the Top	The Shape of Things	S.W.A.T.	Lost in Translation	In the Cut
	budget	2500000	30000000	0	80000000	4000000	12000000
	revenue	12300000	19526014	735992	207700000	119723856	23726793
	runtime	106	87	96	117	102	119
	release_date	2003-03-07 00:00:00	2003-03-21 00:00:00	2003-07-24 00:00:00	2003-08-08 00:00:00	2003-09-18 00:00:00	2003-10-22 00:00:00
	cast_order	3	2	0	5	1	1
	actor_name	Mark Ruffalo	Mark Ruffalo	Paul Rudd	Jeremy Renner	Scarlett Johansson	Mark Ruffalo
	imdb_votes	26,182	29,610	11,962	156,467	507,784	26,960
	Internet Movie Database	7.4	5.3	6.6	6.1	7.7	5.4
	Rotten Tomatoes	NaN	14.0	64.0	48.0	95.0	35.0
	Metacritic	57.0	27.0	59.0	45.0	91.0	47.0
	opening_weekend	40515.0	7600000.0	173246.0	37062535.0	925087.0	97625.0
	Action	0	0	0	1	0	0
	Adventure	0	0	0	0	0	0
	Comedy	0	1	1	0	1	0

Drama	1	1	1	0	1	1
Family	0	0	0	0	0	0
Fantasy	0	0	0	0	0	0
Horror	0	0	0	0	0	0
Music	0	0	0	0	0	0
Mystery	0	0	0	0	0	1
Romance	1	1	1	0	1	1
Science Fiction	0	0	0	0	0	0
Thriller	0	0	0	1	0	1
War	0	0	0	0	0	0
Western	0	0	0	0	0	0
MCU	0	0	0	0	0	0
MCU Entry	0	0	0	0	0	0
release_year	2003	2003	2003	2003	2003	2003
release_quarter	1	1	3	3	3	4
Treatment	1	1	1	1	1	1
Lead Role	0	0	1	0	1	1
Supporting Role	1	1	0	1	0	0
MCU Entry Year	0	0	0	0	0	0
opening_weekend_log	10.609428	15.843659	12.062468	17.428117	13.737643	11.488889

35 rows x 670 columns

```

model2 = treatment_control.groupby(["actor_name","release_year"]).agg(count_lead_
                                average_rat
                                film_count=
                                opening_wee
                                Treatment=('

model2.dropna(axis=0,inplace=True)

```

```
model2.head()
```

	actor_name	release_year	count_lead_roles	average_rating	film_count	open
0	Ben Foster	2003	0	6.65	2	
2	Ben Foster	2005	1	6.50	1	
3	Ben Foster	2006	0	6.75	2	
4	Ben Foster	2007	0	7.10	2	
6	Ben Foster	2009	2	6.90	2	

✓ Exploratory Data Analysis

```
df_event = df_event[df_event['Treatment'] == 1] # only treated actors
df_event['event_time'] = df_event['release_year'] - df_event['MCU Entry Year']
df_event = df_event[df_event['event_time'].between(-5, 10)]

# Group: average outcome by event time
agg = df_event.groupby('event_time')['opening_weekend_log'].mean().reset_index()


pivot = treatment_control.pivot_table(
    index='actor_name',
    columns='release_year',
    values='opening_weekend_log',
    aggfunc='mean'
)

plt.figure(figsize=(14, 10))
sns.heatmap(pivot, cmap='YlGnBu', linewidths=0.1, linecolor='gray', cbar_kws={'label': 'Opening Weekend Log Revenue'})
plt.title("Heatmap: Opening Weekend Log Revenue per Actor-Year")
plt.xlabel("Year")
plt.ylabel("Actor")
plt.tight_layout()
plt.show()
```

```
<ipython-input-224-a3a6793fb099>:2: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/st>



```
treatment_control.reset_index(drop=True).to_json("treatment_control.json",orient=''
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

```
# Load your dataset (adjust path if needed)
df = pd.read_json("treatment_control.json")
```

```
# Convert key numeric columns
df['opening_weekend_log'] = pd.to_numeric(df.get('opening_weekend_log'), errors='')
df['budget'] = pd.to_numeric(df.get('budget'), errors='coerce')
df['revenue'] = pd.to_numeric(df.get('revenue'), errors='coerce')
df['runtime'] = pd.to_numeric(df.get('runtime'), errors='coerce')
df['imdb_votes'] = pd.to_numeric(df.get('imdb_votes'), errors='coerce')
df['release_year'] = pd.to_numeric(df.get('release_year'), errors='coerce')

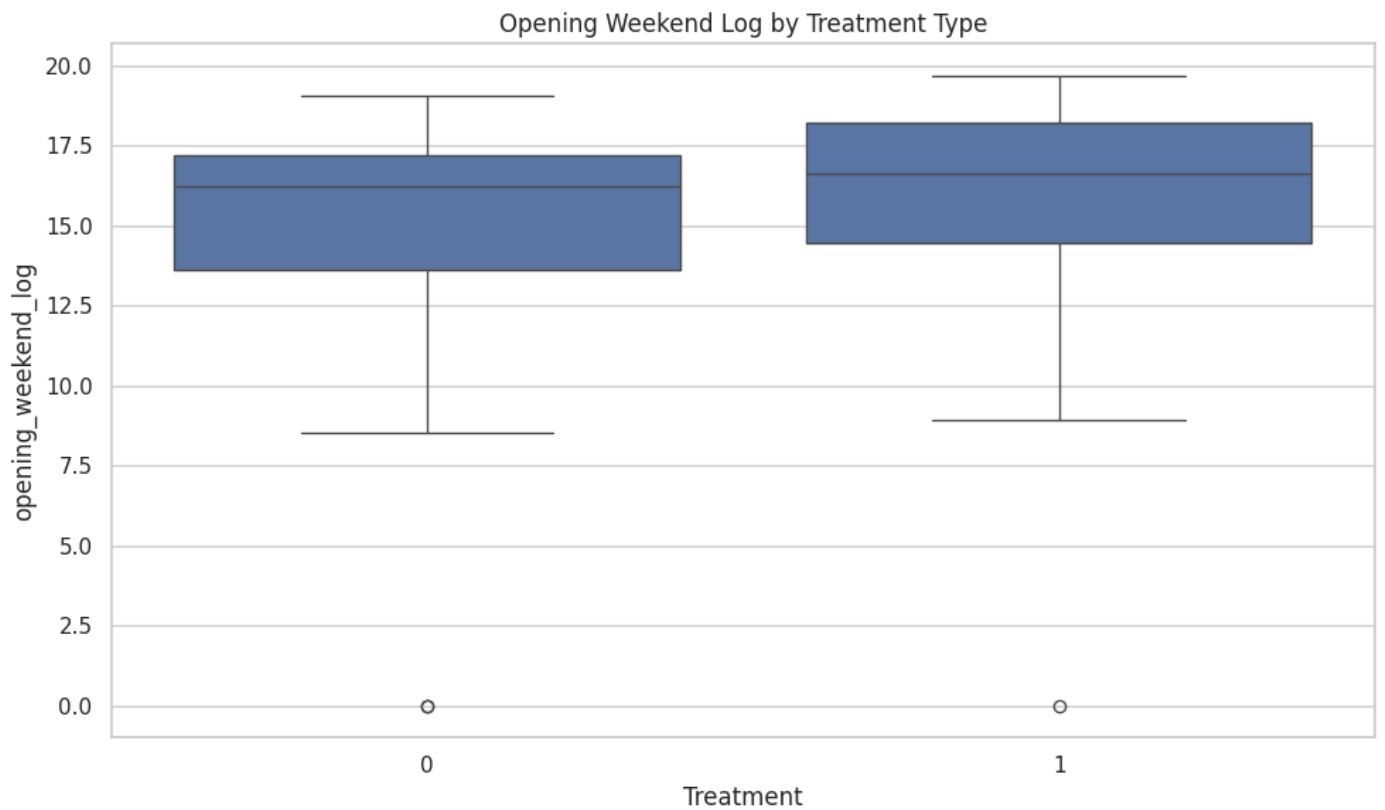
# Set seaborn style
sns.set(style="whitegrid")
```

```
numeric_cols = ['opening_weekend_log', 'budget', 'revenue', 'runtime', 'imdb_vote
```


```
# 2. Correlation Heatmap
plt.figure(figsize=(10, 6))
corr = df[numeric_cols].corr()
sns.heatmap(corr, annot=True, cmap="coolwarm", vmin=-1, vmax=1)
plt.title("Correlation Heatmap")
plt.tight_layout()
plt.show()
```




```
# 3. Boxplot: Treatment vs Opening Weekend  
plt.figure(figsize=(10, 6))  
sns.boxplot(x='Treatment', y='opening_weekend_log', data=df)  
plt.title("Opening Weekend Log by Treatment Type")  
plt.tight_layout()  
plt.show()
```



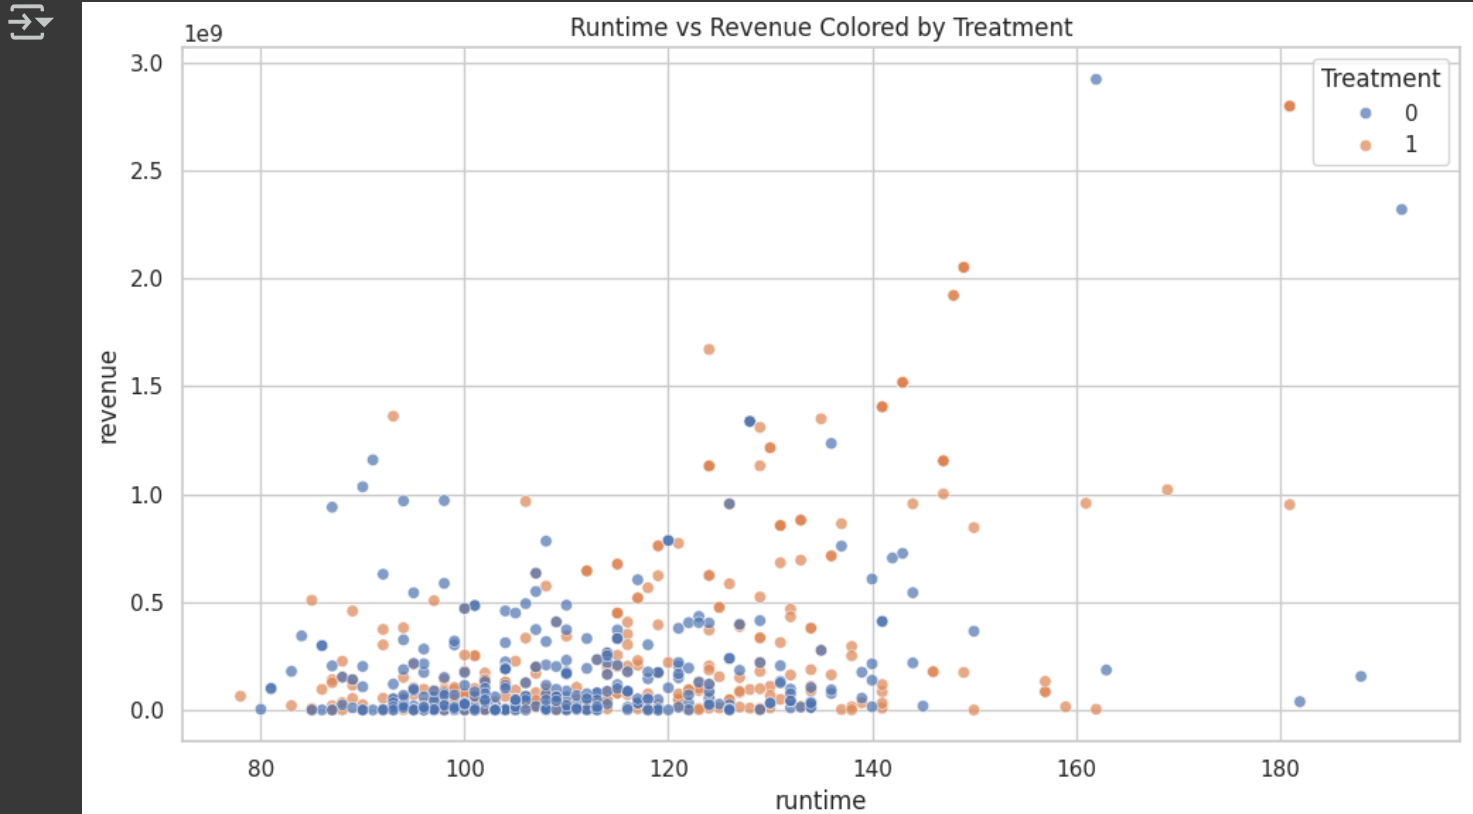
```
# 4. Distribution: Lead Role
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='opening_weekend_log', hue='Lead Role', kde=True, element:
plt.title("Opening Weekend Log Distribution by Lead Role")
plt.tight_layout()
plt.show()
```



```

a# 5. Scatterplot: Runtime vs Revenue by Treatment
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='runtime', y='revenue', hue='Treatment', alpha=0.7)
plt.title("Runtime vs Revenue Colored by Treatment")
plt.tight_layout()
plt.show()

```




```


plt.figure(figsize=(10, 6))
sns.lineplot(data=df, x='release_year', y='opening_weekend_log', hue='Treatment',
plt.title("Pre-Trend Check: Avg. Opening Weekend Log Revenue by Year")
plt.axvline(x=2008, color='black', linestyle='--', label="First MCU Entry")
plt.xlabel("Release Year")

```

```
plt.ylabel("Avg. Log Opening Revenue")
plt.legend(title="Treated (1) vs Control (0)")
plt.grid(True)
plt.tight_layout()
plt.show()
```

 <ipython-input-248-318f76580552>:2: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.




```
plt.figure(figsize=(10, 6))
```

```
plt.figure(figsize=(10, 6))
sns.lineplot(
    data=df,
    x='release_year',
    y='Lead Role',
    hue='Treatment',
    estimator='mean',
    ci=None
)
plt.axvline(x=2008, color='black', linestyle='--', label="First MCU Entry")
plt.title("Lead Role Share Over Time by Treatment Group")
plt.ylabel("Proportion in Lead Role")
plt.grid(True)
plt.tight_layout()
plt.show()
```

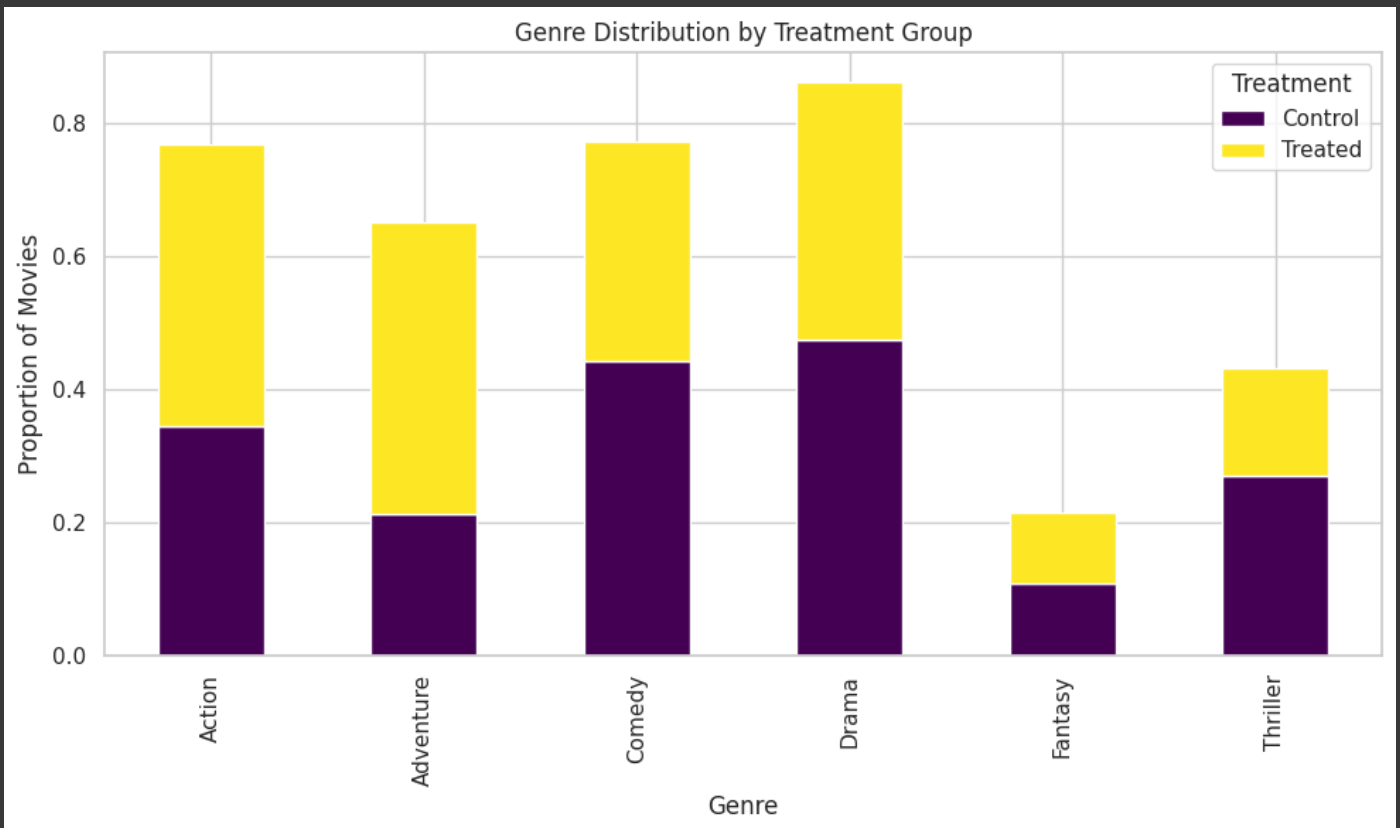
 <ipython-input-250-41f3763148ea>:2: FutureWarning:

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.



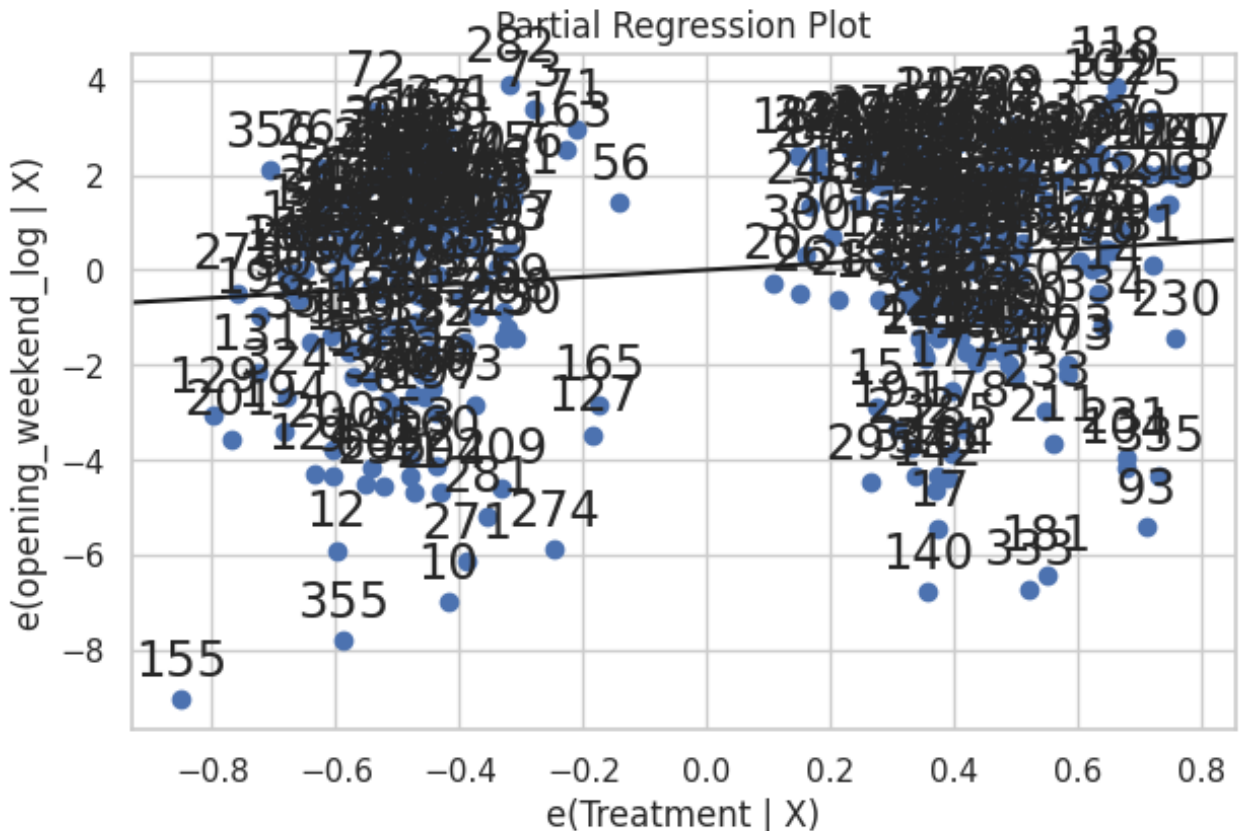
```
genre_cols = ['Action', 'Adventure', 'Comedy', 'Drama', 'Fantasy', 'Thriller']
genre_share = df.groupby('Treatment')[genre_cols].mean().T

genre_share.plot(kind='bar', stacked=True, figsize=(10, 6), colormap='viridis')
plt.title("Genre Distribution by Treatment Group")
plt.ylabel("Proportion of Movies")
plt.xlabel("Genre")
plt.legend(title="Treatment", labels=['Control', 'Treated'])
plt.tight_layout()
plt.show()
```



```
# Run OLS model with relevant controls
model = smf.ols(
    formula='opening_weekend_log ~ Treatment + average_rating + film_count + count',
    data=model2
).fit()

# Plot partial regression (added variable plot)
fig = sm.graphics.plot_partregress('opening_weekend_log', 'Treatment', ['average_rating', 'film_count', 'count'])
fig.suptitle("Partial Regression Plot: Treatment Effect Controlling for Covariate")
plt.tight_layout()
plt.show()
```



```
from sklearn.linear_model import LogisticRegression


# Define feature set for matching
```




```
X = model2[['average_rating', 'film_count', 'count_lead_roles']].dropna()
y = model2.loc[X.index, 'Treatment']


# Fit logistic model
ps_model = LogisticRegression()
ps_model.fit(X, y)
model2.loc[X.index, 'propensity_score'] = ps_model.predict_proba(X)[:, 1]

# Plot
plt.figure(figsize=(10, 6))
sns.histplot(data=model2, x='propensity_score', hue='Treatment', element='step',
plt.title("Propensity Score Distribution by Treatment Group")
plt.tight_layout()
plt.show()
```

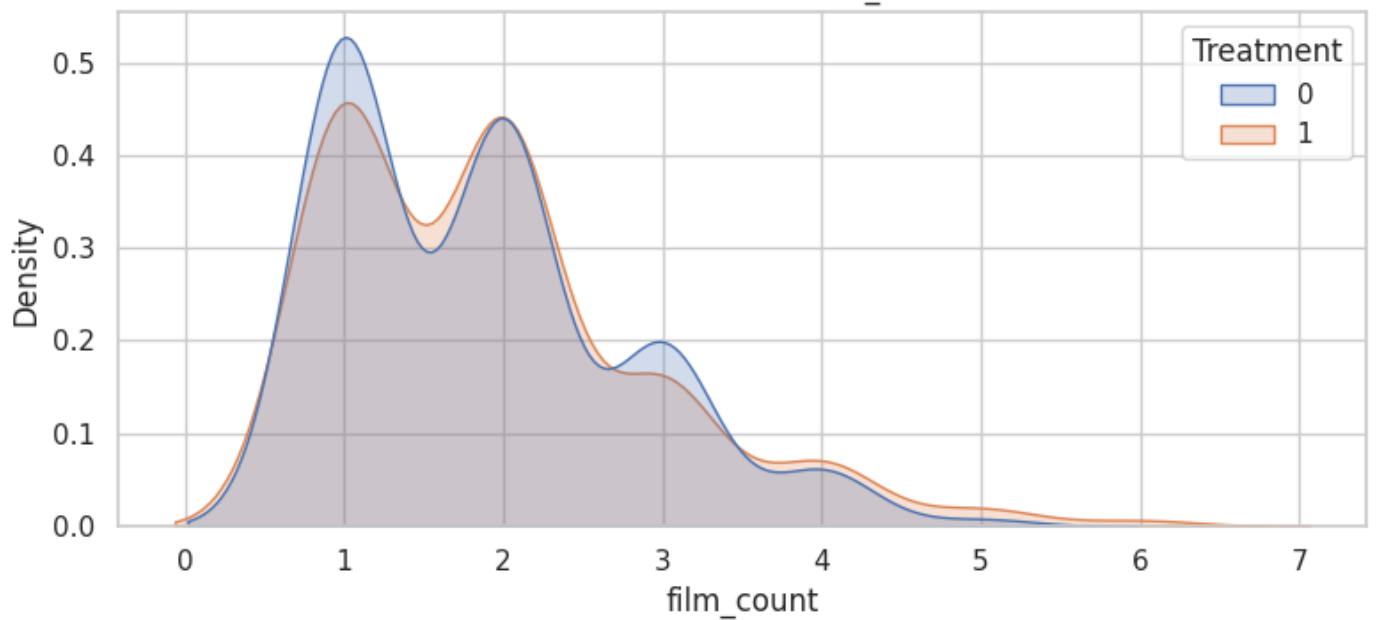


```
# Visualize distribution of covariates
for var in ['average_rating', 'film_count', 'count_lead_roles']:
    plt.figure(figsize=(8, 4))
    sns.kdeplot(data=model2, x=var, hue='Treatment', fill=True, common_norm=False)
    plt.title(f"Covariate Balance: {var}")
    plt.tight_layout()
    plt.show()
```






Covariate Balance: film_count



Covariate Balance: count_lead_roles



`count_lead_roles`

✓ Weighted SMF model

```
# Add inverse probability weights (IPW)
model2['ipw'] = np.where(
    model2['Treatment'] == 1,
    1 / model2['propensity_score'],
    1 / (1 - model2['propensity_score'])
)

# Weighted regression
weighted_model = smf.wls(
    formula='opening_weekend_log ~ Treatment',
    data=model2,
    weights=model2['ipw']
).fit()

print(weighted_model.summary())
```



WLS Regression Results

```
=====
Dep. Variable:      opening_weekend_log    R-squared:                0.001
Model:              WLS                   Adj. R-squared:           0.001
Method:             Least Squares         F-statistic:             10.8
Date:               Wed, 07 May 2025       Prob (F-statistic):      0.0016
Time:               18:00:04              Log-Likelihood:          -799.8
No. Observations:   346                   AIC:                     1604
Df Residuals:       344                   BIC:                     1611
Df Model:           1
Covariance Type:    nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	15.1462	0.184	82.364	0.000	14.784	15.508
Treatment	0.8504	0.258	3.296	0.001	0.343	1.358

```
=====
Omnibus:              82.047    Durbin-Watson:           1.732
Prob(Omnibus):        0.000    Jarque-Bera (JB):       156.772
Skew:                 -1.274    Prob(JB):               9.07e-35
Kurtosis:              5.094    Cond. No.                2.64
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correct

✓ Difference-in-Differences (DiD) Model Estimation

```
# 2. Estimate the DiD with controls + fixed effects
```

```
# 3. Estimate the DID with controls + fixed effects
formula = (
    'opening_weekend_log ~ Treatment'
    ' + film_count + average_rating + count_lead_roles'
    ' + C(actor_name) + C(release_year)'
)

did_model = smf.ols(formula, data=model2) \
    .fit(cov_type='cluster', cov_kws={'groups': model2['actor_name']})

print(did_model.summary())
```



OLS Regression Results

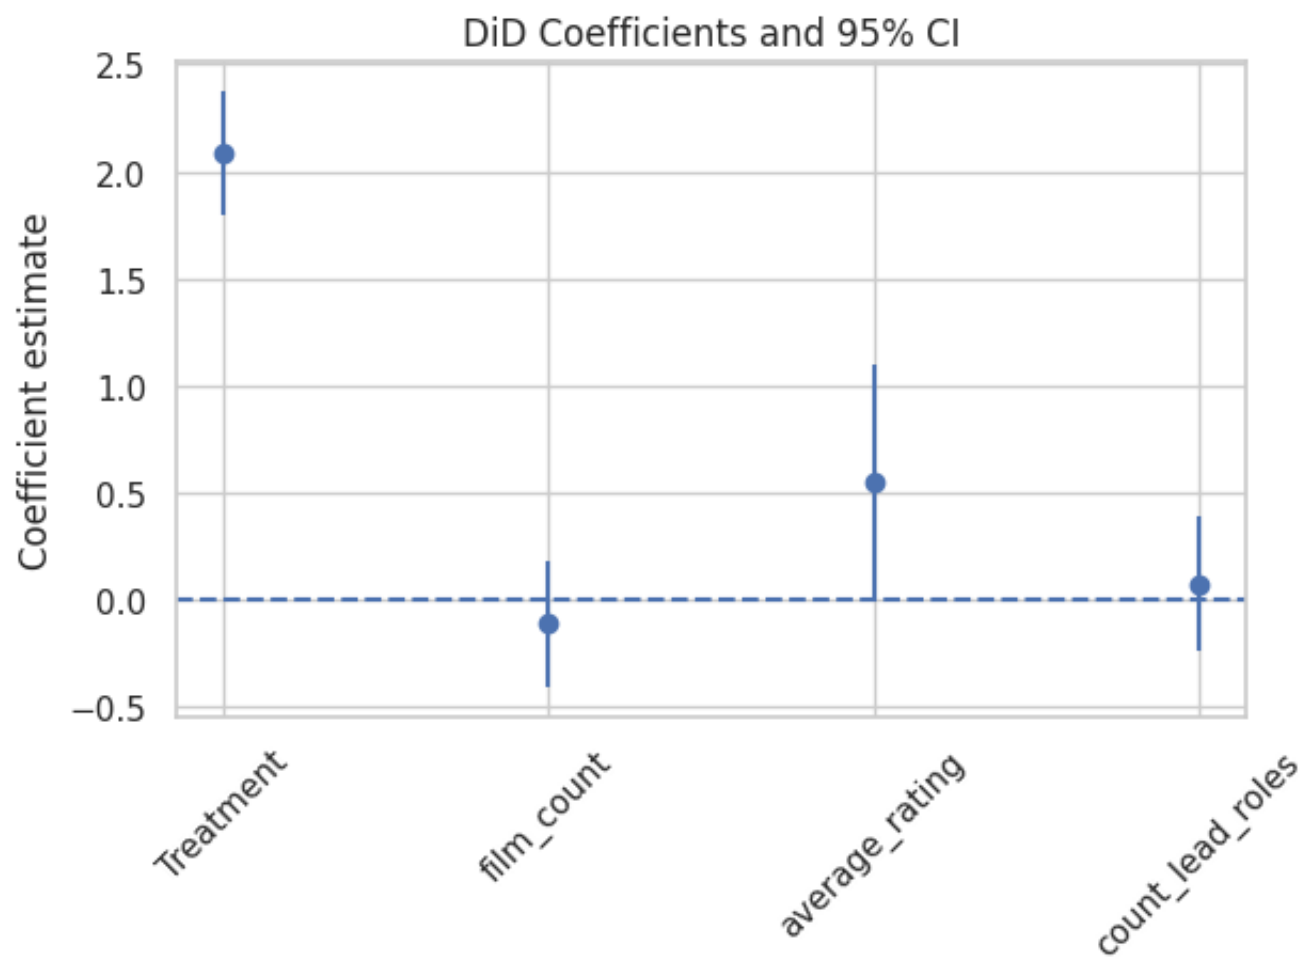
```
=====
Dep. Variable:      opening_weekend_log    R-squared:                0.35
Model:              OLS                   Adj. R-squared:           0.25
Method:             Least Squares         F-statistic:             107.1
Date:               Wed, 07 May 2025      Prob (F-statistic):      2.60e-1
Time:               18:05:40              Log-Likelihood:          -718.5
No. Observations:   346                  AIC:                    1533
Df Residuals:       298                  BIC:                    1718
Df Model:           47
Covariance Type:    cluster
=====
```

	coef	std err	z	P
Intercept	9.4373	1.999	4.722	(
C(actor_name)[T.Benedict Cumberbatch]	-0.2466	0.095	-2.596	(
C(actor_name)[T.Chadwick Boseman]	0.6344	0.157	4.042	(
C(actor_name)[T.Channing Tatum]	2.5380	0.135	18.792	(
C(actor_name)[T.Charlize Theron]	1.9890	0.149	13.372	(
C(actor_name)[T.Chris Evans]	0.7595	0.078	9.711	(
C(actor_name)[T.Chris Hemsworth]	0.2444	0.126	1.939	(
C(actor_name)[T.Chris Pratt]	0.7973	0.114	7.004	(
C(actor_name)[T.Edward Norton]	0.8664	0.165	5.253	(
C(actor_name)[T.Jeremy Renner]	-0.1319	0.085	-1.546	(
C(actor_name)[T.Jim Carrey]	2.4628	0.198	12.439	(
C(actor_name)[T.John David Washington]	1.2770	0.339	3.769	(
C(actor_name)[T.Mark Ruffalo]	-0.4173	0.094	-4.422	(
C(actor_name)[T.Matthew Goode]	0.3049	0.174	1.752	(
C(actor_name)[T.Paul Rudd]	0.5115	0.096	5.339	(
C(actor_name)[T.Robert Downey Jr.]	0.8082	0.115	7.024	(
C(actor_name)[T.Ryan Reynolds]	2.2230	0.182	12.206	(
C(actor_name)[T.Sam Worthington]	1.0058	0.178	5.655	(
C(actor_name)[T.Scarlett Johansson]	0.1605	0.090	1.784	(
C(actor_name)[T.Steve Carell]	2.4944	0.146	17.045	(
C(actor_name)[T.Taron Egerton]	1.7221	0.277	6.213	(
C(actor_name)[T.Tom Hiddleston]	-0.5398	0.156	-3.466	(
C(actor_name)[T.Tom Holland]	-0.4950	0.121	-4.095	(

C(actor_name) [T.Tye Sheridan]	-0.1266	0.199	-0.636	(
C(release_year) [T.2004]	-1.1611	1.346	-0.862	(
C(release_year) [T.2005]	0.7634	0.790	0.966	(
C(release_year) [T.2006]	-0.9767	1.233	-0.792	(
C(release_year) [T.2007]	-0.6335	0.981	-0.646	(
C(release_year) [T.2008]	0.3939	0.893	0.441	(
C(release_year) [T.2009]	1.2495	0.922	1.356	(
C(release_year) [T.2010]	-0.2065	1.275	-0.162	(
C(release_year) [T.2011]	-0.0612	0.846	-0.072	(
C(release_year) [T.2012]	1.3094	0.932	1.405	(
C(release_year) [T.2013]	0.2939	0.970	0.303	(
C(release_year) [T.2014]	-0.6417	0.976	-0.658	(
C(release_year) [T.2015]	0.7931	1.029	0.771	(
C(release_year) [T.2016]	1.8709	0.991	1.889	(
C(release_year) [T.2017]	0.7094	1.135	0.625	(
C(release_year) [T.2018]	2.0673	0.997	2.073	(
C(release_year) [T.2019]	2.2219	0.896	2.481	(
C(release_year) [T.2020]	1.3749	0.871	1.579	(
C(release_year) [T.2021]	1.6262	0.991	1.642	(
C(release_year) [T.2022]	2.6143	0.834	3.134	(
C(release_year) [T.2023]	0.8428	1.137	0.741	(
C(release_year) [T.2024]	1.6230	0.792	2.050	(

```
# 3. Pull out the four key coefficients + CIs
vars_to_plot = ['Treatment', 'film_count', 'average_rating', 'count_lead_roles'] #
params = did_model.params[vars_to_plot]
ci_low, ci_high = did_model.conf_int().loc[vars_to_plot].T.values

# 4. Draw the coefficient plot
plt.figure()
plt.errorbar(
    vars_to_plot,
    params.values,
    yerr=[params.values - ci_low, ci_high - params.values],
    fmt='o'
)
plt.axhline(0, linestyle='--')
plt.xticks(rotation=45)
plt.ylabel('Coefficient estimate')
plt.title('DiD Coefficients and 95% CI')
plt.tight_layout()
plt.show()
```



```
df_event = treatment_control.copy()
```